


Visualizing ASH

John Beresniewicz
NoCOUG 2018



Agenda

- What is ASH?
 - Mechanism and properties
 - Usage: ASH Math, Average Active Sessions
- ASH Visualizations
 - EM Performance: Wait class details, Top Activity,
 - EM Treemaps: Enterprise Loadmap, ASH Analytics
- Visual Investigation of ASH dumps using R and ggplot

What is ASH?

- Active Session History (V\$ACTIVE_SESSION_HISTORY)
- Session State Objects sampled every 1000 ms by MMNL
- ACTIVE (non-idle) session state data captured into circular memory buffer
- Latch-less and efficient, sampling is independent of session activity
- “Fix-up” mechanism updates recent samples with new data (TIME_WAITED not known when sampled session is waiting)


Standard usage: ASH Math

- **COUNT(*) = DB Time (seconds)**
- ASH is a FACT table where each row equals one second of session DB Time
- Use GROUP BY to break down DB Time by any of the many dimensions of ASH
- **Aggregate ASH samples to Decompose DB Time over various dimension combinations**


Visualizing ASH in EM

- Performance Page > Wait Class drill-down
- Top Activity Page
- ASH Analytics
- SQL Monitor


EM Performance Page



EM Top Activity - ASH





Average Active Sessions



ASH Treemaps

[refresh](#)

Hierarchy: SQLID > WAITCLASS > INSTANCE_NAME



visible depth

2



smallest area value

1.0



largest area value

141.0



smallest color value

-0.04



largest color value

0.02



[refresh](#)

Hierarchy: INSTANCE > WAITCLASS > EVENT



bug1ap
User I/O
db file sequential read

Cluster
read by others
CPU
ON CPU
db file scatter
Commit
log file
cold

bug2ap
User I/O
db file sequential read
db file scat
db file pr

Other
cr request retry
CPU
ON CPU
Cluster
gc buffer
gc c
g g
gc cr rec
g g
g g
Commit
Cor
log file
curs
lib
row
Sys

visible depth


smallest area value

largest area value

smallest color value

largest color value

EM ASH Analytics



Exploring ASH dumps visually using R and ggplot

Tools and raw data

- Tools: Anaconda, R, RStudio, ggplot2, dplyr
(Hadley Wickham's stuff)
- Data: some old 10g ASH dumps
 - 10.2.0.1 Enterprise Edition
 - 4-node RAC ASH dump > 4 separate trace files

Data import and preparation

- Getting started was easy:

```
ashDF <- read.csv(ashfilepath, header=FALSE)
```


- HOWEVER: real analysis required significant data munging
 - multiple trace files per dump
 - join in wait event and wait class names
 - import microsecond sample times
 - add state boolean: time_waited > 0 = WAIT else CPU
 - add “CPU” wait-class

DB Time over Time (AAS)


What and why?

- Standard usage:
Aggregate ASH samples into Average Active Sessions
- Investigate consistency across RAC instances
- Investigate sampler independence, sample id consistency?
- Compute AAS over wait classes by 1-minute intervals


$x = \text{sample time, color wait class}$




color by instance



color wait class, facet instance




$x = \text{sample id, color wait class}$




color by instance, hmmm




facet instance, color wait class







sample id vs. sample time




AAS by minute, color wait class



Wait class colors



It took significant effort to finalize the wait class color scheme.




ASH Sampler timing


What and why?

- Micro-second sample times in ASH dump data
- Does the sampler do well at keeping to fixed interval?
- Use inter-sample time diffs to analyze sampler consistency
- `diff_msecs` = `sample_time - lag(sample_time)`
expressed in microseconds, grouped by instance number


density plot of diff_msecs




violin plot of diff_msecs




plot diff msec - 1000




color by instance




filter “good” ones out




skewed distribution with long tail




zoom in on the action




diffs in 10 ms intervals?



ha! (serialization, scheduler?)




consistency across samplers




See what we got:
observed event counts


Box plot of $\log_{10}(\text{TIME_WAITED})$




latency box plots by wait class




by event id




not helpful



faceted by wait class



free x-axis so geoms fill space




See what could be:
estimated event counts


What and why?

- Patent [US9633061B2](#) granted in 2017
(Uri Shaft, Graham Wood & John Beresniewicz)
- Methods for determining event counts based on time-sampled data
(i.e. ASH)
- Estimate average event latencies from samples:
 - est latency = est DBtime / est count
 - **EST_COUNT = MAX(1000000 / TIME_WAITED, 1)**
(when $0 < \text{time_waited}$, and sample interval = 1000 ms)
- Performance Metric: Estimated Event Latency over past minute


density plot: observed event count by latency




latency density weighted by estimated count




together
(blue = observed, red = estimated)




violin density plot, observed




faceted by instance




weighted by estimated count




together, the balloon squeeze




looks cool, but is it useful?



facet by wait class, misleading




much better view




Back to raw data


all samples over time - overplots




reduce point size and alpha




facet wait class, theme_dark




theme_bw



reduced alpha and point size



estimated 1 min avg latencies
color by instance, facet by wait class



ASH Visualization

- ASH data is highly dimensional and temporally fine-grained
- Ask questions and visualize to investigate
- Understand the visualization and whether it answers question
- Iterate visualizations by changing geometry attributes
- Experiment liberally, but be true to the data